Chahat Raj

PhD in Computer Science | George Mason University

chahatraj.github.io craj@gmu.edu github.com/chahatraj Google Scholar

Research Interests

I work on Responsible AI, Ethics, and Fairness, specifically in the evaluation and mitigation of socio-cultural biases within multilingual and multimodal NLP applications with a focus on LLMs. I have recently published at EMNLP, AIES, and ECIR.

Education

Present – Aug 2022 | **George Mason University** — **Fairfax, USA**
Ph.D. in Computer Science
Advisors: Ziwei Zhu, Antonios Anastasopoulos

Aug 2021 – Aug 2019 | **Delhi Technological University** — **Delhi, India**
Masters in Information Systems (Research Track)
Advisor: Priyanka Meel

May 2019 – Aug 2015 | **Indira Gandhi Delhi Technical University for Women** — **Delhi, India**
Bachelors in Computer Science & Engineering

Experience

Sep 2024 – May 2024 | **University of Washington | The Information School** — **Seattle, WA**
Researcher | Advisor: Aylin Caliskan
Project: Social Biases in Visual Question Answering in Generative AI.

Aug 2024 – May 2024 | **George Mason University | School of Computer Science** — **Fairfax, VA**
Graduate Research Assistant | Advisors: Ziwei Zhu, Antonios Anastasopoulos
Project: Exploring Stereotypical Associations in Large Vision-Language Models through generative tasks - image generation and question answering.

Aug 2023 – May 2023 | **George Mason University | School of Computer Science** — **Fairfax, VA**
Graduate Research Assistant | Advisor: Ziwei Zhu
Project: Fairness and Interpretability of Transformer Models.

Jul 2022 – Jun 2021 | **University of Technology Sydney | School of Computer Science** — **Remote**
Visiting Scholar | Advisor: Mukesh Prasad
Projects: Disaster Fundraising on Social Media Analysis with the Australian Red Cross, AI integration in aboriginal and indigenous communities' lifestyle, Web information pollution detection (cyberbullying, hate, and offensive speech), Machine learning-based decision-making in clinical vertigo diagnosis.

Aug 2021 – May 2021 | **Indian Institute of Management Raipur** — **Remote**
Researcher | Advisor: Manojit Chattopadhyay
Projects: Cryptocurrencies' price prediction, Analysing causal relationships between stock market behavior and societal events.

Aug 2022 – Aug 2019 | **Delhi Technological University** — **Delhi, India**
Graduate Student Researcher | Advisor: Priyanka Meel
Projects: Developing neural network frameworks to detect antisocial web content (fake news, infodemic, misinformation, disinformation, rumor, death hoax, and clickbait), Infodemic impact analysis.

Selected Publications

[C.5] **BiasDora: Exploring Hidden Biased Associations in Vision-Language Models** [PDF | Code]
Chahat Raj, Anjishnu Mukherjee, Aylin Caliskan, Antonios Anastasopoulos, Ziwei Zhu
Conference on Empirical Methods in Natural Language Processing **[EMNLP'24]**

[C.4] **Breaking Bias, Building Bridges: Evaluation and Mitigation of Social Biases in LLMs via Contact Hypothesis** [PDF | Code]
Chahat Raj, Anjishnu Mukherjee, Aylin Caliskan, Antonios Anastasopoulos, Ziwei Zhu
AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society **[AIES'24]**

[C.3] **SALSA: Salience-Based Switching Attack for Adversarial Perturbations in Fake News Detection Models** [PDF | Code]
Chahat Raj*, Anjishnu Mukherjee*, Hemant Purohit, Antonios Anastasopoulos, Ziwei Zhu
European Conference on Information Retrieval **[ECIR'24]**

[C.2] **Global Voices, Local Biases: Socio-cultural Prejudices across Languages** [PDF | Code]
Anjishnu Mukherjee*, Chahat Raj*, Ziwei Zhu, Antonios Anastasopoulos
Conference on Empirical Methods in Natural Language Processing **[EMNLP'23]**

[C.1] **True and Fair: Robust and Unbiased Fake News Detection via Interpretable Machine Learning** [PDF | Code]
Chahat Raj, Anjishnu Mukherjee, Ziwei Zhu
AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society **[AIES'23]**

Skills

Natural Language Processing, Generative AI, Multimodal Large Language Models, Model Training, Evaluation & Validation, Instruction-tuning, Prompt Engineering, LLM-human alignment, Vision, Ethics and Bias Mitigation, Explainability.

Talks

"Breaking Bias, Building Bridges"
› Talk at AAAI/ACM AIES — October 2024 (San Jose, USA)

"LLM Ethics - Bias and Mitigation"
› Invited Lecture - CS 678 Advanced NLP @ George Mason University — October 2024 (Virginia, USA)

"A Psychological View to Social Biases in LLMs"
› SouthNLP @ Emory University — April 2024 (Georgia, USA)

Awards

Travel Grant by AAAI/ACM AIES 2024 & 2023 Funded $1000 to travel to San Jose, USA and $1000 to travel to Montreal, Canada to attend the AIES conference.

CAHMP Fellowship @ George Mason University, 2024 Summer research on social biases in question answering.

Best Paper Award @ MASC-SLL 2024, Johns Hopkins University For the paper 'A Psychological View to Social Bias in LLMs: Evaluation and Mitigation'.

Summer Graduate Research Award @ George Mason University, 2023 For research on 'Explainability and Robustness of Debiasing Approaches for Transformers'.

Research Excellence Award @ Delhi Technological University 2023 $(\times 2)$ **& 2022** Received three Commendable Research Awards with a total cash prize of INR 150k by DTU, Delhi, for the research published in reputed scientific journals.

Graduate Scholarship @ Delhi Technological University 2019 Received INR 297k by AICTE for qualifying GATE.

Teaching

CS 108 Introduction to Computer Programming, GMU *Teaching Assistant* — Spring'24
› Responsible for teaching labs, one-to-one student tutoring, creating lab modules, grading, and exam invigilation.

CS 478 Natural Language Processing, GMU *Teaching Assistant* — Fall'23
› Delivered ad-hoc lectures, student tutoring, teaching labs, and assignment grading.

COMP 502 Mathematical Foundations of Computing, GMU *Teaching Assistant* — Fall'23
› Responsible for one-to-one student tutoring sessions, and assignment grading.

CS 112 Introduction to Python Programming, GMU *Teaching Assistant* — Fall'22, Spring'23
› Responsible for teaching labs, one-to-one student tutoring, creating lab modules, assignment grading, and exam invigilation.

Academic Service

Reviewer ACM Transactions on Intelligent Systems and Technology 2024, LTEDI @ EACL 2024, TrustNLP 2024, CSCW 2024, NeurIPS 2023, Elsevier & Springer Journal papers
Sub-Reviewer EMNLP'24, AIES'24